

# Terry Stops Classification

Student name: Cassidy Exum Student pace: self paced Scheduled project review date/time: Undetermined Instructor name: Morgan Jones Blog post URL: <https://exumexaminesdata.blogspot.com/2022/10/the-chess-cheating-scandal-and-72-page.html>

## Objective - Create a classification model that can predict whether an arrest was made or not made

The data for this project was obtained from the seattle .gov website at the follow link: <https://data.seattle.gov/Public-Safety/Terry-Stops/28ny-9ts8>

The stakeholder is a non-profit organization researching police stop data.

Description of problem taken from Flatiron phase 3 dataset github located: <https://github.com/learn-co-curriculum/dsc-phase-3-choosing-a-dataset>

In Terry v. Ohio, a landmark Supreme Court case in 1967-8, the court found that a police officer was not in violation of the "unreasonable search and seizure" clause of the Fourth Amendment, even though he stopped and frisked a couple of suspects only because their behavior was suspicious. Thus was born the notion of "reasonable suspicion", according to which an agent of the police may e.g. temporarily detain a person, even in the absence of clearer evidence that would be required for full-blown arrests etc. Terry Stops are stops made of suspicious drivers.

Build a classifier to predict whether an arrest was made after a Terry Stop, given information about the presence of weapons, the time of day of the call, etc. This is a binary classification problem.

Note that this dataset also includes information about gender and race. You may use this data as well. You could conceivably pitch your project as an inquiry into whether race (of officer or of subject) plays a role in whether or not an arrest is made.

If you do elect to make use of race or gender data, be aware that this can make your project a highly sensitive one; your discretion will be important, as well as your transparency about how you use the data and the ethical issues surrounding it.

Brief idea of how I should do this -

1. fix data
2. get dummies

3. train model (Use Logistic as Baseline model)
4. iterate (xgboost/random forest -> GridSearchCV xgboost/random forest)

## Phase 1 - Imports and Data exploration

Import the data set, explore, clean, preprocess, and reorganize the data.

```
In [1]: #relevant imports

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

# sklearn imports
# Model Selection and Preprocessing
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.model_selection import cross_val_score

# Metrics
from sklearn.metrics import accuracy_score, f1_score, recall_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_curve, auc

# Classifiers
from xgboost import XGBClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
```

```
In [2]: #import data and print head

df = pd.read_csv('terry_stops.csv')
df.head()
```

```
Out[2]:
```

	Subject Age Group	Subject ID	GO / SC Num	Terry Stop ID	Stop Resolution	Weapon Type	Officer ID	Officer YOB	Offic Gend
0	-	-8	20220000063036	32023419019	Field Contact	-	6805	1973	
1	-	-8	20220000233868	35877423282	Field Contact	-	8881	1988	
2	-	-1	20140000120677	92317	Arrest	None	7500	1984	
3	-	-1	20150000001463	28806	Field Contact	None	5670	1965	

Subject Age Group	Subject ID	GO / SC Num	Terry Stop ID	Stop Resolution	Weapon Type	Officer ID	Officer YOB	Offic Gend
4	-	-1	20150000001516	29599	Field Contact	None	4844	1961

5 rows × 23 columns

Preliminary thoughts

There are a few ID type columns I don't need such as Subject ID, GO/ SC Num , Terry Stop ID.

There are a few officer specific columns which may or may not be useful such as Officer ID, YOB, officer squad

Location Info probably not needed - Precinct, Sector, Beat

Target = Stop Resolution

```
In [3]: #check entire dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53654 entries, 0 to 53653
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Subject Age Group                    53654 non-null  object
1   Subject ID                          53654 non-null  int64
2   GO / SC Num                        53654 non-null  int64
3   Terry Stop ID                      53654 non-null  int64
4   Stop Resolution                     53654 non-null  object
5   Weapon Type                        53654 non-null  object
6   Officer ID                         53654 non-null  object
7   Officer YOB                        53654 non-null  int64
8   Officer Gender                     53654 non-null  object
9   Officer Race                       53654 non-null  object
10  Subject Perceived Race              53654 non-null  object
11  Subject Perceived Gender            53654 non-null  object
12  Reported Date                      53654 non-null  object
13  Reported Time                      53654 non-null  object
14  Initial Call Type                  53654 non-null  object
15  Final Call Type                    53654 non-null  object
16  Call Type                          53654 non-null  object
17  Officer Squad                      53165 non-null  object
18  Arrest Flag                        53654 non-null  object
19  Frisk Flag                         53654 non-null  object
20  Precinct                          53654 non-null  object
21  Sector                            53654 non-null  object
22  Beat                              53654 non-null  object
dtypes: int64(4), object(19)
memory usage: 9.4+ MB
```

```
In [4]: # Create Month and Weak features from Reported Date
df['Reported Date'] = pd.to_datetime(df['Reported Date'])
df['Month'] = df['Reported Date'].apply(lambda x: x.month)
```

```
In [5]: #Create a function to map the day to a week 1-4. Create new feature, week of the
def week_map(x):
```

```

if x.day<=7:
    return 1
elif x.day<=14:
    return 2
elif x.day<=21:
    return 3
else:
    return 4

```

```
df['Week'] = df['Reported Date'].apply(week_map)
```

```

In [6]: # Convert Reported Time to binary night or day
df['day/night'] = df['Reported Time'].apply(lambda x: 'night' if '00:00' <= x <=

```

```

In [7]: # ID columns, overlapping columns, and anything else
to_drop = ['Subject ID',
           'GO / SC Num',
           'Terry Stop ID',
           'Officer ID',
           'Reported Time',
           'Initial Call Type',
           'Final Call Type',
           'Officer Squad']
df.drop(to_drop, axis=1, inplace=True)
df.head()

```

```

Out[7]:

```

	Subject Age Group	Stop Resolution	Weapon Type	Officer YOB	Officer Gender	Officer Race	Subject Perceived Race	Subject Perceived Gender	Reported Date
0	-	Field Contact	-	1973	M	White	DUPLICATE	DUPLICATE	2022-03-00:00:00+00:
1	-	Field Contact	-	1988	M	Asian	DUPLICATE	DUPLICATE	2022-09-00:00:00+00:
2	-	Arrest	None	1984	M	Black or African American	Asian	Male	2015-10-00:00:00+00:
3	-	Field Contact	None	1965	M	White	-	-	2015-03-00:00:00+00:
4	-	Field Contact	None	1961	M	White	White	Male	2015-03-00:00:00+00:

```

In [8]: #check target variable
df['Stop Resolution'].value_counts()

```

```

Out[8]:
Field Contact      23006
Offense Report     16599
Arrest             13131
Referred for Prosecution    728
Citation / Infraction    190
Name: Stop Resolution, dtype: int64

```

I can reorganize this column to only be 'Arrest' and 'No Arrest' or 1 and 0.

```

In [9]: # Change everything other than Arrest to 0, and Arrest to 1
df.loc[ df['Stop Resolution'] != 'Arrest', 'Stop Resolution'] = 0

```

```
df.loc[ df['Stop Resolution'] == 'Arrest', 'Stop Resolution'] = 1
df['Stop Resolution'].value_counts()
```

```
Out[9]: 0    40523
        1    13131
        Name: Stop Resolution, dtype: int64
```

```
In [10]: # Set type
df['Stop Resolution'] = df['Stop Resolution'].astype('int64')
```

```
In [11]: #Check Call Type
df['Call Type'].value_counts()
```

```
Out[11]: 911                24803
        -                13514
        ONVIEW            11184
        TELEPHONE OTHER, NOT 911    3685
        ALARM CALL (NOT POLICE ALARM)  446
        TEXT MESSAGE           21
        SCHEDULED EVENT (RECURRING)    1
        Name: Call Type, dtype: int64
```

```
In [12]: #Check Precinct column
df['Precinct'].value_counts()
```

```
Out[12]: West            14070
        North           11699
        -              10240
        East            6904
        South           6363
        Southwest       2320
        SouthWest       1775
        Unknown         200
        OOJ             61
        FK ERROR        22
        Name: Precinct, dtype: int64
```

```
In [13]: # Impute - as unknown some values
        # Impute FK Error as Unknown
        # Fix Southwest
df.loc[ df['Precinct'] == '-', 'Precinct'] = 'Unknown'
df.loc[ df['Precinct'] == 'FK ERROR', 'Precinct'] = 'Unknown'
df.loc[ df['Precinct'] == 'SouthWest', 'Precinct'] = 'Southwest'
```

```
In [14]: #Check sector column
df['Sector'].value_counts()
```

```
Out[14]: -    10477
        K     4411
        M     4369
        E     3589
        N     3219
        D     3094
        B     2524
        F     2487
        R     2353
        L     2335
        Q     2192
        O     2031
        S     1980
        U     1963
        G     1742
```

```
J      1657
W      1606
C      1572
99      53
Name: Sector, dtype: int64
```

```
In [15]: # impute 99 to be -
df.loc[ df['Sector'] == '99', 'Sector'] = '-'
```

```
In [16]: #Check Beat Column
df['Beat'].value_counts()
```

```
Out[16]: -      10385
K3       2374
M3       1950
N3       1608
E2       1529
M2       1215
M1       1207
D1       1202
N2       1187
D2       1166
E1       1153
K2       1100
R2       1076
Q3       1045
F2       1017
B1        941
K1        937
B2        934
U2        922
E3        906
O1        844
L2        798
L1        790
S2        788
F3        756
L3        747
D3        726
F1        714
R1        697
W2        680
U1        662
Q2        658
S3        657
B3        651
G2        640
J3        631
O3        626
C1        618
J1        612
G3        607
R3        580
O2        561
C3        553
W1        537
S1        535
G1        494
Q1        489
N1        423
J2        414
C2        402
W3        390
U3        379
```

```
99          100
OOJ          39
S            2
Name: Beat, dtype: int64
```

```
In [17]: # impute 99 as -
df.loc[ df['Beat'] == '99', 'Beat'] = '-'
```

```
In [18]: # Check the column
df['Subject Perceived Race'].value_counts()
```

```
Out[18]:
```

White	26320
Black or African American	15936
Unknown	3526
-	1810
Asian	1803
Hispanic	1684
American Indian or Alaska Native	1514
Multi-Racial	809
Other	152
Native Hawaiian or Other Pacific Islander	98
DUPLICATE	2

Name: Subject Perceived Race, dtype: int64

```
In [19]: # drop DUPLICATE and impute - as Unknown
df = df[df['Subject Perceived Race'] != 'DUPLICATE']
```

```
In [20]: # Impute - as unknown
df.loc[ df['Subject Perceived Race'] == '-', 'Subject Perceived Race'] = 'Unknown'
df['Subject Perceived Race'].value counts()
```

```
Out[20]:
```

White	26320
Black or African American	15936
Unknown	5336
Asian	1803
Hispanic	1684
American Indian or Alaska Native	1514
Multi-Racial	809
Other	152
Native Hawaiian or Other Pacific Islander	98

Name: Subject Perceived Race, dtype: int64

```
In [21]: #Check counts of subject gender
df['Subject Perceived Gender'].value counts()
```

```
Out[21]: Male 42251
Female 10749
Unable to Determine 326
- 239
Unknown 67
Gender Diverse (gender non-conforming and/or transgender) 20
Name: Subject Perceived Gender, dtype: int64
```

```
In [22]: #impute - and Unable to Determine as Unknown
df.loc[ df['Subject Perceived Gender'] == '-', 'Subject Perceived Gender'] = 'Un
df.loc[ df['Subject Perceived Gender'] == 'Unable to Determine', 'Subject Percei
df['Subject Perceived Gender'].value counts()
```

```
Out[22]: Male          42251
         Female        10749
         Unknown         632
```

Gender Diverse (gender non-conforming and/or transgender)  
Name: Subject Perceived Gender, dtype: int64

20

```
In [23]: #check weapon type column
df['Weapon Type'].value_counts()
```

```
Out[23]: None                    32565
-                    17798
Lethal Cutting Instrument      1482
Knife/Cutting/Stabbing Instrument  967
Handgun                        342
Blunt Object/Striking Implement  125
Firearm Other                  100
Firearm                        63
Club, Blackjack, Brass Knuckles  49
Mace/Pepper Spray              44
Other Firearm                  41
Firearm (unk type)             15
Taser/Stun Gun                 13
Fire/Incendiary Device         11
None/Not Applicable            10
Club                           9
Rifle                          8
Shotgun                        4
Personal Weapons (hands, feet, etc.)  2
Automatic Handgun              2
Brass Knuckles                 1
Blackjack                      1
Name: Weapon Type, dtype: int64
```

```
In [24]: # Impute - as None
df.loc[ df['Weapon Type'] == '-', 'Weapon Type'] = 'None'
df['Weapon Type'].value_counts()
```

```
Out[24]: None                    50363
Lethal Cutting Instrument      1482
Knife/Cutting/Stabbing Instrument  967
Handgun                        342
Blunt Object/Striking Implement  125
Firearm Other                  100
Firearm                        63
Club, Blackjack, Brass Knuckles  49
Mace/Pepper Spray              44
Other Firearm                  41
Firearm (unk type)             15
Taser/Stun Gun                 13
Fire/Incendiary Device         11
None/Not Applicable            10
Club                           9
Rifle                          8
Shotgun                        4
Personal Weapons (hands, feet, etc.)  2
Automatic Handgun              2
Brass Knuckles                 1
Blackjack                      1
Name: Weapon Type, dtype: int64
```

```
In [25]: # Simplify Weapon Type Further
# Organize into Blunt Weapons, Sharp Weapons, Firearms
#Firearms
df.loc[ df['Weapon Type'] == 'None/Not Applicable', 'Weapon Type'] = 'None'
df.loc[ df['Weapon Type'] == 'Handgun', 'Weapon Type'] = 'Firearm'
df.loc[ df['Weapon Type'] == 'Firearm Other', 'Weapon Type'] = 'Firearm'
df.loc[ df['Weapon Type'] == 'Other Firearm', 'Weapon Type'] = 'Firearm'
```



```

df.loc[ df['Weapon Type'] == 'Firearm (unk type)', 'Weapon Type'] = 'Firearm'
df.loc[ df['Weapon Type'] == 'Rifle', 'Weapon Type'] = 'Firearm'
df.loc[ df['Weapon Type'] == 'Shotgun', 'Weapon Type'] = 'Firearm'
df.loc[ df['Weapon Type'] == 'Automatic Handgun', 'Weapon Type'] = 'Firearm'

# Blunt Object/Striking Implement
df.loc[ df['Weapon Type'] == 'Blackjack', 'Weapon Type'] = 'Blunt Object/Strikin
df.loc[ df['Weapon Type'] == 'Brass Knuckles', 'Weapon Type'] = 'Blunt Object/St
df.loc[ df['Weapon Type'] == 'Club', 'Weapon Type'] = 'Blunt Object/Striking Imp
df.loc[ df['Weapon Type'] == 'Club, Blackjack, Brass Knuckles', 'Weapon Type'] =
df.loc[ df['Weapon Type'] == 'Personal Weapons (hands, feet, etc.)', 'Weapon Typ

#knives
df.loc[ df['Weapon Type'] == 'Knife/Cutting/Stabbing Instrument', 'Weapon Type']

#Other
df.loc[ df['Weapon Type'] == 'Mace/Pepper Spray', 'Weapon Type'] = 'Other'
df.loc[ df['Weapon Type'] == 'Taser/Stun Gun', 'Weapon Type'] = 'Other'
df.loc[ df['Weapon Type'] == 'Fire/Incendiary Device', 'Weapon Type'] = 'Other'

```

```

In [26]: #check weapon type
df['Weapon Type'].value_counts()

```

```

Out[26]: None                    50373
Lethal Cutting Instrument      2449
Firearm                        575
Blunt Object/Striking Implement  187
Other                          68
Name: Weapon Type, dtype: int64

```

```

In [27]: # Check officer gender
df['Officer Gender'].value_counts()

```

```

Out[27]: M      47514
F         6108
N           30
Name: Officer Gender, dtype: int64

```

```

In [28]: # Check officer YOB
df['Officer YOB'].value_counts()

```

```

Out[28]: 1986      3690
1987      3422
1991      2979
1984      2921
1992      2854
1990      2688
1985      2600
1988      2395
1989      2272
1982      1946
1983      1866
1993      1776
1995      1716
1979      1715
1981      1591
1994      1346
1971      1272
1976      1246
1978      1221
1977      1101
1973      1004

```

```

1996      962
1980      935
1967      792
1997      746
1970      670
1968      664
1969      590
1975      579
1974      579
1962      463
1964      459
1972      449
1965      424
1963      265
1966      235
1961      234
1958      222
1959      174
1960      161
1998      123
1900       69
1954       44
1957       43
1953       35
1999       25
2000       23
1955       21
1956       17
1948       11
1952        9
1949        5
1946        2
1951        1
Name: Officer YOB, dtype: int64

```

```

In [29]: #yob as type int
df['Officer YOB'].astype('int64')

```

```

Out[29]: 2      1984
3      1965
4      1961
5      1963
6      1977
...
53649   1977
53650   1996
53651   1973
53652   1978
53653   1995
Name: Officer YOB, Length: 53652, dtype: int64

```

```

In [30]: #drop people born in 1900, 122 year olds arent still officers
df = df[df['Officer YOB'] != 1900]

# turn YOB to age
df['Officer Age Group'] = df['Officer YOB'].apply(
    lambda x: 2022-x)

```

```

In [31]: # Map age to age brackets
def map_age(x):
    if 18 <= x <= 25:
        return '18 - 25'
    elif 26 <= x <= 35:

```

```

        return '26 - 35'
    elif 36 <= x <= 45:
        return '36 - 45'
    elif 46 <= x <= 55:
        return '46 - 55'
    elif 56 <= x:
        return '56 and Above'

```

```
df['Officer Age Group'] = df['Officer Age Group'].apply(map_age)
```

```
In [32]: #check that the map worked
df['Officer Age Group'].value_counts()
```

```
Out[32]: 26 - 35          22410
36 - 45          19586
46 - 55           7845
56 and Above     2825
18 - 25           917
Name: Officer Age Group, dtype: int64
```

```
In [33]: #drop Officer YOB
df.drop('Officer YOB', axis=1, inplace=True)
```

```
In [34]: #check officer Race
df['Officer Race'].value_counts()
```

```
Out[34]: White          39375
Two or More Races      3336
Hispanic or Latino     3278
Asian                  2398
Not Specified          2293
Black or African American 2098
Nat Hawaiian/Oth Pac Islander 472
American Indian/Alaska Native 333
Name: Officer Race, dtype: int64
```

```
In [35]: #Make Not Specified as Unknown
df.loc[ df['Officer Race'] == 'Not Specified', 'Officer Race'] = 'Unknown'
```

```
In [36]: #Check age group column
df['Subject Age Group'].value_counts()
```

```
Out[36]: 26 - 35          17912
36 - 45          11618
18 - 25          10499
46 - 55           6882
56 and Above     2779
1 - 17           2079
-                1814
Name: Subject Age Group, dtype: int64
```

```
In [37]: # check arrest flag column
df['Arrest Flag'].value_counts()
```

```
Out[37]: N      48668
Y       4915
Name: Arrest Flag, dtype: int64
```

```
In [38]: #check frisk flag
df['Frisk Flag'].value_counts()
```

```
Out[38]: N    40736
        Y    12369
        -      478
        Name: Frisk Flag, dtype: int64
```

```
In [39]: #final check
        df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 53583 entries, 2 to 53653
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Subject Age Group                    53583 non-null  object
1   Stop Resolution                      53583 non-null  int64
2   Weapon Type                         53583 non-null  object
3   Officer Gender                      53583 non-null  object
4   Officer Race                        53583 non-null  object
5   Subject Perceived Race              53583 non-null  object
6   Subject Perceived Gender            53583 non-null  object
7   Reported Date                       53583 non-null  datetime64[ns, UTC]
8   Call Type                           53583 non-null  object
9   Arrest Flag                         53583 non-null  object
10  Frisk Flag                           53583 non-null  object
11  Precinct                             53583 non-null  object
12  Sector                              53583 non-null  object
13  Beat                                53583 non-null  object
14  Month                               53583 non-null  int64
15  Week                                53583 non-null  int64
16  day/night                           53583 non-null  object
17  Officer Age Group                   53583 non-null  object
dtypes: datetime64[ns, UTC](1), int64(3), object(14)
memory usage: 7.8+ MB
```

```
In [40]: #No longer need reported Date
        df = df.drop(columns='Reported Date')
```

```
In [41]: #I think this field contradicts the Stop Resolution field, and based on the colu
        #Stop Resolution is more official so we will drop this one.

        #create a copy of the df right before dropping this to compare later -
        df_with_arrest_flag = df.copy()
        df = df.drop(columns='Arrest Flag')
```

## One Hot Encode the data

```
In [42]: #obtain dummies
        dummy_df = pd.get_dummies(df, drop_first=True)
```

## Phase 2 - Modeling

```
In [43]: # Data X and Y

        y = dummy_df['Stop Resolution']
        X = dummy_df.drop('Stop Resolution', axis=1)
```

```
In [44]: # Split the data
        X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=123, test
```

```
In [45]: # Baseline model

logistic_regression_clf = LogisticRegression(random_state=123)

logistic_regression_clf.fit(X_train, y_train)

/Users/cassidyexum/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
Out[45]: LogisticRegression(random_state=123)
```

```
In [46]: #predictions and test score

train_preds = logistic_regression_clf.predict(X_train)
test_preds = logistic_regression_clf.predict(X_test)

logistic_regression_clf.score(X_test, y_test)
```

```
Out[46]: 0.7587339504329651
```

```
In [47]: print(classification_report(y_test, test_preds))
```

	precision	recall	f1-score	support
0	0.77	0.97	0.86	10170
1	0.49	0.08	0.14	3226
accuracy			0.76	13396
macro avg	0.63	0.53	0.50	13396
weighted avg	0.70	0.76	0.69	13396

```
In [48]: print(confusion_matrix(y_test, test_preds))

[[9907  263]
 [2969  257]]
```

```
In [49]: cross_val_score(logistic_regression_clf, X_train, y_train, cv=3)
```

/Users/cassidyexum/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/sklearn/linear\_model/\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.  
  
 Increase the number of iterations (max\_iter) or scale the data as shown in:  
 https://scikit-learn.org/stable/modules/preprocessing.html  
 Please also refer to the documentation for alternative solver options:  
 https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression  
 n\_iter\_i = \_check\_optimize\_result(  
 /Users/cassidyexum/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/sklearn/linear\_model/\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n
n_iter_i = _check_optimize_result(
/Users/cassidyexum/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n
n_iter_i = _check_optimize_result(
```

Out[49]: array([0.7539564 , 0.75335921, 0.75483389])

The learner is about 75% accurate, could be random guessing no arrest, or it just fits this the general proportions of the data. Need to go more in depth.

## Perceptron

```
In [50]: from sklearn.linear_model import Perceptron

#Instantiate
perceptron_clf = Perceptron()

#fit
perceptron_clf.fit(X_train, y_train)

print('Perceptron Test Score: ', perceptron_clf.score(X_test, y_test))
test_preds = perceptron_clf.predict(X_test)
train_preds = perceptron_clf.predict(X_train)
print('=====')

#classification report
print('Test Classification Report')
print(classification_report(y_test, test_preds))
print('=====')
print('Train Classification Report')
print(classification_report(y_train, train_preds))
```

```
Perceptron Test Score:  0.7393998208420424
=====
Test Classification Report
```

	precision	recall	f1-score	support
0	0.77	0.93	0.84	10170
1	0.38	0.13	0.19	3226
accuracy			0.74	13396
macro avg	0.58	0.53	0.52	13396
weighted avg	0.68	0.74	0.69	13396

```
=====
Train Classification Report
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.76	0.93	0.84	30295
1	0.38	0.12	0.18	9892
accuracy			0.73	40187
macro avg	0.57	0.53	0.51	40187
weighted avg	0.67	0.73	0.68	40187

## Decision Tree with GridSearchCV

```
In [51]: # instantiate DT classifier and obtain cross val score
dt_clf = DecisionTreeClassifier()
dt_cv_score = cross_val_score(dt_clf, X_train, y_train, cv=3)
mean_dt_cv_score = np.mean(dt_cv_score)

print(f"Mean Cross Validation Score: {mean_dt_cv_score :.2%}")
```

Mean Cross Validation Score: 68.64%

```
In [52]: #params for grid search
dt_param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 2, 4, 6],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 3, 5]
}
```

```
In [53]: # instantiate grid search and fit
dt_grid_search = GridSearchCV(dt_clf, dt_param_grid, cv=3, return_train_score=True)

# Fit to the data
dt_grid_search.fit(X_train, y_train)

# Obtain the parameters of the best tree
print('Best Params: ', dt_grid_search.best_params_)

#best score check, is it better than before?
print('Best Score: ', dt_grid_search.best_score_)
```

Best Params: {'criterion': 'gini', 'max\_depth': 4, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2}  
Best Score: 0.7541244839066019

```
In [54]: # Instantiate classifier with best params
dt_clf = DecisionTreeClassifier(criterion='gini',
                                max_depth=4,
                                min_samples_leaf=1,
                                min_samples_split=2)
```

```
In [55]: #Fit classifier
dt_clf.fit(X_train, y_train)
```

Out[55]: DecisionTreeClassifier(max\_depth=4)

```
In [56]: # Obtain preds and scores
test_preds = dt_clf.predict(X_test)
train_preds = dt_clf.predict(X_train)
```

```
print('Train Score: ', dt_clf.score(X_train, y_train))
print('=====')
print('Test Score: ', dt_clf.score(X_test, y_test))
```

```
Train Score:  0.7542737701246672
=====
Test Score:   0.7595550910719617
```

```
In [57]: # Classification report
print(classification_report(y_test, test_preds))
```

	precision	recall	f1-score	support
0	0.76	1.00	0.86	10170
1	0.55	0.01	0.02	3226
accuracy			0.76	13396
macro avg	0.66	0.50	0.44	13396
weighted avg	0.71	0.76	0.66	13396

```
In [58]: print(classification_report(y_train, train_preds))
```

	precision	recall	f1-score	support
0	0.75	1.00	0.86	30295
1	0.57	0.01	0.01	9892
accuracy			0.75	40187
macro avg	0.66	0.50	0.44	40187
weighted avg	0.71	0.75	0.65	40187

Obtaining poor results, essentially the same scores as before. Hopefully KNN gets a better results or we move to XGBoost for better results. If nothing works, I have to go back to the beginning and keep more columns. It will increase runtime but give the learners more to learn from.

## XGBoost attempt -

KNN seems to be taking way too long to run so I'm going to bypass it and move to XGBoost.

As a recap - Logistic Regression was about 73% accurate and a Decision Tree using Grid Search got us to 75%. XGBoost is considered one of the best so I'm hoping to get to 80%

```
In [59]: #Instantiate the classifier
xgb_clf = XGBClassifier()

# fit the classifier
xgb_clf.fit(X_train, y_train)
```

```
Out[59]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.300000012, max_delta_step=0, max_depth=6,
                        min_child_weight=1, missing=nan, monotone_constraints='()',
                        n_estimators=100, n_jobs=0, num_parallel_tree=1, random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=None)
```





```

min_child_weight=1, missing=nan,
monotone_constraints='()',
n_estimators=100, n_jobs=0,
num_parallel_tree=1, random_state=0,
reg_alpha=0, reg_lambda=1,
scale_pos_weight=1, subsample=1,
tree_method='exact', validate_parameters=1,
verbosity=None),
param_grid={'booster': ['gbtree'],
            'learning_rate': [0.25, 0.3, 0.35],
            'max_depth': [5, 6, 7],
            'n_estimators': [75, 100, 125]},
return_train_score=True)

```

```

In [64]: #obtain best params and best score
print('Best Score: ', xgb_grid_search.best_score_)
print("=====")
print('Best Params: ', xgb_grid_search.best_params_)

```

Best Score: 0.7520591852113654

=====

Best Params: {'booster': 'gbtree', 'learning\_rate': 0.25, 'max\_depth': 5, 'n\_estimators': 100}

## Re doing the base XGBoost model (best predictor) and obtaining coefficients

```

In [65]: #Instantiate the classifier
xgb_clf = XGBClassifier()

# fit the classifier
xgb_clf.fit(X_train, y_train)

#Train Scores
train_preds = xgb_clf.predict(X_train)

#Test Scores
test_preds = xgb_clf.predict(X_test)

#classification report
print('Test Classification Report')
print(classification_report(y_test, test_preds))
print('=====')
print('Train Classification Report')
print(classification_report(y_train, train_preds))
print('=====')
print('Train Score: ', xgb_clf.score(X_train, y_train))
print('=====')
print('Test Score: ', xgb_clf.score(X_test, y_test))
print('=====')

```

Test Classification Report

	precision	recall	f1-score	support
0	0.78	0.95	0.86	10170
1	0.48	0.14	0.22	3226
accuracy			0.76	13396
macro avg	0.63	0.55	0.54	13396
weighted avg	0.71	0.76	0.70	13396

```

=====
Train Classification Report
      precision    recall  f1-score   support

     0       0.80      0.98      0.88      30295
     1       0.79      0.25      0.38       9892

 accuracy         0.80      40187
 macro avg       0.80      0.61      0.63      40187
weighted avg       0.80      0.80      0.76      40187

=====
Train Score:  0.799213676064399
=====
Test Score:   0.7566437742609734
=====

```

```

In [66]: feature_importance = pd.DataFrame(xgb_clf.feature_importances_)
feature_importance = feature_importance.T
feature_importance.columns = X_train.columns

```

```

In [67]: print(feature_importance.columns)

```

```

Index(['Month', 'Week', 'Subject Age Group_1 - 17',
      'Subject Age Group_18 - 25', 'Subject Age Group_26 - 35',
      'Subject Age Group_36 - 45', 'Subject Age Group_46 - 55',
      'Subject Age Group_56 and Above', 'Weapon Type_Firearm',
      'Weapon Type_Lethal Cutting Instrument',
      ...,
      'Beat_U2', 'Beat_U3', 'Beat_W1', 'Beat_W2', 'Beat_W3',
      'day/night_night', 'Officer Age Group_26 - 35',
      'Officer Age Group_36 - 45', 'Officer Age Group_46 - 55',
      'Officer Age Group_56 and Above'],
      dtype='object', length=120)

```

```

In [68]: feature_importance.T.sort_values(0, ascending=False)

```

```

Out[68]:

```

	0
Precinct_Unknown	0.425884
Frisk Flag_Y	0.017746
Precinct_North	0.017104
Call Type_911	0.013095
Beat_M2	0.011080
...	...
Subject Perceived Race_Other	0.001238
Subject Perceived Gender_Gender Diverse (gender non-conforming and/or transgender)	0.001220
Beat_OOJ	0.000333
Beat_S	0.000000
Call Type_SCHEDULED EVENT (RECURRING)	0.000000

120 rows x 1 columns

```
In [69]: # Consolidate dummy importance values into a single value for the original column

importance_consolidated = pd.DataFrame(np.zeros(df.columns.size))
importance_consolidated = importance_consolidated.T
importance_consolidated.columns = df.columns
```

```
for con_col in importance_consolidated.columns:
    for feat_col in feature_importance.columns:
        if con_col in feat_col:
            importance_consolidated[con_col] += feature_importance[feat_col]

importance_consolidated.drop('Stop Resolution', axis=1, inplace=True)
```

```
In [70]: importance_consolidated
```


```
Out[70]:
```

	Subject Age Group	Weapon Type	Officer Gender	Officer Race	Subject Perceived Race	Subject Perceived Gender	Call Type	Frisk Flag	Precinct	Stop Resolution
0	0.02606	0.022233	0.005031	0.031379	0.033484	0.008795	0.032855	0.02767	0.468378	0.000000

```
In [71]: values = importance_consolidated.T
values = values[0].sort_values(0, ascending=False)
```

```
In [72]: ticks = importance_consolidated.T
ticks = ticks[0].sort_values(0, ascending=False).index
```

```
In [73]: fig = plt.figure(figsize = (16, 8))
plt.bar(x = ticks,
        height = values)
plt.xticks(rotation=60)
plt.xlabel('Features')
plt.ylabel('Feature Importance')
plt.title('Features Sorted by Importance');
```



```
In [74]: feature_importance
```

Out[74]:

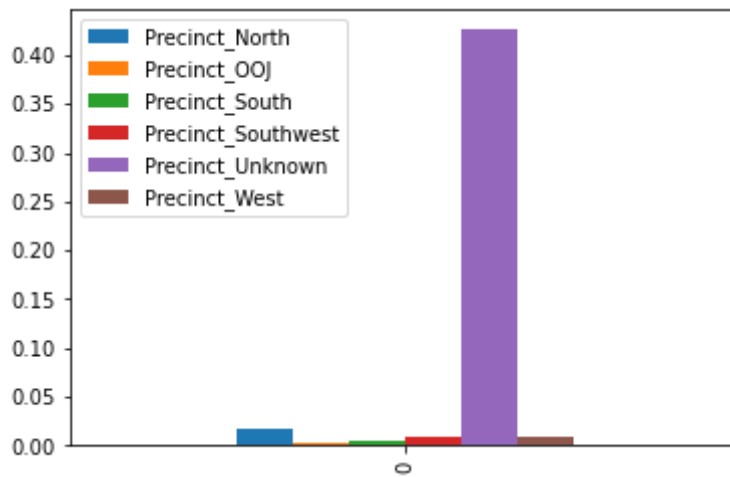
	Month	Week	Subject Age Group_1 - 17	Subject Age Group_18 - 25	Subject Age Group_26 - 35	Subject Age Group_36 - 45	Subject Age Group_46 - 55	Subject Age Group_56 and Above	W Type_Fi
0	0.003988	0.003823	0.004788	0.00427	0.004034	0.003999	0.004572	0.004397	0.0

1 rows x 120 columns

```
In [75]: precinct_list = []
for col in feature_importance:
    if 'Precinct' in col:
        precinct_list.append(col)

precinct_df = feature_importance[precinct_list]
precinct_df.plot.bar()
```

Out[75]: <AxesSubplot:>



## Business Recommendations and conversation

The problem statement for this was to create a predictor. I created 4, and all 4 had scores of around 75-80%

In the context of the data, this information should be used strictly for research purposes. In no way should a machine learning model be used to determine whether or not to arrest someone.

My recommendation would be to investigate why precinct is the largest predictor. If we had socioeconomic data we might be able to draw more conclusions about different areas of Seattle and why they see more arrests than others. I would call on the city to provide census data to allow for this further research.

## Testing - Create 50/50 split dataset and determine if the model is working or if just guessing No Arrest.

In [ ]:

## Testing - Adding back in Arrest Flag and seeing results

In [76]:

```
#df_with_arrest_flag
new_dummies = pd.get_dummies(df_with_arrest_flag, drop_first=True)
```

In [77]:

```
y = new_dummies['Stop Resolution']
X = new_dummies.drop('Stop Resolution', axis=1)

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=123, test
```

In [78]:

```
#Instantiate the classifier
xgb_clf_2 = XGBClassifier()

# fit the classifier
xgb_clf_2.fit(X_train, y_train)

#Train Scores
```

```

train_preds = xgb_clf_2.predict(X_train)

#Test Scores
test_preds = xgb_clf_2.predict(X_test)

#classification report
print('Test Classification Report')
print(classification_report(y_test, test_preds))
print('=====')
print('Train Classification Report')
print(classification_report(y_train, train_preds))
print('=====')
print('Train Score: ', xgb_clf_2.score(X_train, y_train))
print('=====')
print('Test Score: ', xgb_clf_2.score(X_test, y_test))
print('=====')

```

```

Test Classification Report
              precision    recall  f1-score   support

     0       0.84         0.98         0.91       10170
     1       0.89         0.43         0.58        3226

 accuracy                   0.85       13396
 macro avg       0.87         0.71         0.75       13396
 weighted avg    0.86         0.85         0.83       13396

```

```

=====
Train Classification Report
              precision    recall  f1-score   support

     0       0.85         0.99         0.92       30295
     1       0.96         0.48         0.64        9892

 accuracy                   0.87       40187
 macro avg       0.91         0.74         0.78       40187
 weighted avg    0.88         0.87         0.85       40187

```

```

=====
Train Score:  0.8665488839674521
=====

```

```

Test Score:  0.8503284562555987
=====

```

Adding the 'Arrest Flag' column improved results by 6-10% between the training and test sets.

If the goal is to use old data where we always have this column available, it is the best predictor

If the goal is to predict the outcome, where this value isn't known, the models perform just barely better than random guessing.

```

In [79]: feature_importance_2 = pd.DataFrame(xgb_clf_2.feature_importances_)
feature_importance_2 = feature_importance_2.T
feature_importance_2.columns = X_train.columns

# Consolidate dummy importance values into a single value for the original column

importance_consolidated_2 = pd.DataFrame(np.zeros(df_with_arrest_flag.columns.size))
importance_consolidated_2 = importance_consolidated_2.T
importance_consolidated_2.columns = df_with_arrest_flag.columns

for con_col in importance_consolidated_2.columns:

```

```

for feat_col in feature_importance_2.columns:
    if con_col in feat_col:
        importance_consolidated_2[con_col] += feature_importance_2[feat_col]

importance_consolidated_2.drop('Stop Resolution', axis=1, inplace=True)

values = importance_consolidated_2.T
values = values[0].sort_values(0, ascending=False)

ticks = importance_consolidated_2.T
ticks = ticks[0].sort_values(0, ascending=False).index

fig = plt.figure(figsize = (16, 8))
plt.bar(x = ticks,
        height = values)
plt.xticks(rotation=60)
plt.xlabel('Features')
plt.ylabel('Feature Importance')
plt.title('Features Sorted by Importance');

```

